المؤتمر الأول للبحوث الطلابية لجامعة الإسماعيلية الجديدة الأهلية
دور الجامعة في التنمية المستدامة والمجتمعية
The Role of The University in Sustainable and Community Development

FIRE-EDGE: A Real-Time Firefighting Edge System with Thermal Imaging, Sensor Fusion, and Augmented Reality

Authors: *Ahmed S. Lotfy, Adham T. Sayed, Mena R. Ghanem*

Affiliation: *New Ismailia National University, Ismailia, Egypt*

Emails: *ahmedsheriflotfy06@gmail.com ,adhamt864@gmail.com, menaghanim863@gmail.com,*

Abstract

Firefighters face life-threatening challenges in smoke-filled environments, with disorientation contributing to 28% of fatalities. This paper presents FIRE-EDGE, a low-cost augmented reality (AR) system integrating edge AI and multi-sensor fusion. Leveraging YOLOv8 for fire detection and a Kalman filter for navigation, the system achieves 85% mAP@0.5 accuracy with 35ms latency on an NVIDIA Jetson Xavier NX. Simulations in ROS/Gazebo demonstrate a 2.16% improvement over baseline models, offering a sustainable solution for firefighter safety.

Keywords: Firefighter Safety, Edge AI, Sensor Fusion, Thermal Imaging, Sustainable Technology

1. Introduction:

- Firefighters face life-threatening disorientation in smoke-filled environments, contributing to 28% of on-duty fatalities globally. Existing augmented reality (AR) systems rely on cloud processing, introducing latency (>500ms) and cost barriers (>$10,000). This study addresses these gaps by proposing a sustainable, edge-AI-powered AR solution to enhance situational awareness and safety.
- Edge-AI Processing: Real-time hazard detection on a Jetson Xavier NX.
- Sensor Fusion: Combines thermal, LiDAR, and IMU data for drift-resistant navigation.
- Cost Efficiency: Total hardware cost <$2,500.

2. Related Work

2.1 AR Systems for Firefighting

Prior studies have explored AR for firefighting but face critical limitations:

- Qwake C-THRU: A commercial AR helmet using thermal imaging but lacks AI-driven hazard detection.
- Zhang et al. (2021): Proposed a cloud-based AR system with 92% accuracy but suffered from 800ms latency.
- Lee et al. (2022): Developed an IMU-based navigation system with 0.5m/min drift but ignored thermal data.

2.2 Edge AI in Emergency Response

Edge computing has gained traction for low-latency applications:

- Gupta et al. (2020): Deployed YOLOv5 on drones for wildfire detection but used non-thermal datasets.
- Chen et al. (2023): Achieved 82% mAP@0.5 for fire detection on Raspberry Pi but lacked sensor fusion.

2.3 Sensor Fusion Techniques

- Patel et al. (2019): Fused LiDAR and IMU for indoor navigation but excluded thermal imaging.
- Kim et al. (2021): Used Kalman filtering for drone navigation but required cloud post-processing.

2.4 Gaps Addressed by FIRE-EDGE:

- Eliminates cloud dependency.
- Integrates thermal imaging with edge AI.
- Optimized for cost and power efficiency.

3. Methodology

3.1 Hardware Design

Component	Model	Technical Specification	Rationale
NVIDIA Jetson Xavier NX	Jetson Xavier NX 16GB	- GPU: 384-core Volta™ (48 Tensor Cores) - CPU: 6-core Carmel ARM® v8.2 - RAM: 16GB LPDDR4x - Power: 10W (15W max)	Optimized for edge-AI with 21 TOPS performance, enabling real-time YOLOv8 inference at the edge. Provides necessary processing power for object detection and sensor fusion without cloud dependency. Power efficient for battery operation.
FLIR Lepton 3.5	FLIR Lepton 3.5 (Radiometric)	- Resolution: 160×120 - FOV: 57° × 44° - Spectral Range: 8–14 μm - Power: 150mW	Low-cost, compact thermal imaging for fire/human detection in smoke. Allows "seeing" through smoke by detecting heat signatures. Small size and low power consumption suitable for a wearable AR system.
LiDAR	Garmin LIDAR-Lite v3HP	- Range: 40m - Accuracy: ±2.5cm - Power: 1.8W	High-precision obstacle mapping for navigation in low-visibility conditions. Creates a 3D map of the environment, enabling firefighters to avoid obstacles and navigate safely.
IMU	TDK InvenSense MPU-9250	- Gyroscope: ±2000$°/s$ - Accelerometer: ±16g - Magnetometer: ±4800μT - Power: 3.3V, 3.5mA	9-DoF (Degrees of Freedom) Inertial Measurement Unit (IMU) for motion tracking and dead reckoning. Tracks orientation and movement to maintain accurate position tracking when GPS is unavailable (e.g., indoors).
Ultrasonic Sensors	A02YYUW Waterproof	- Range: 20cm–4.5m - Accuracy: ±1cm - Power: 3.3V–5V, 3mA	Waterproof design for harsh environments; lower power consumption with high precision.
Custom Battery Pack	Custom 4S LiPo	- Voltage: 14.8V - Capacity: 10,000mAh - Lifespan: 45 minutes at 12W load	Balances portability and runtime for field operations.
Display	Kopin Lightning OLED (0.5")	- Resolution: 1280×720 (HD) - Size: 0.5" diagonal - Brightness: 10,000 nits - Power: 1.5W	Ultra-compact, high-brightness display for AR overlays in smoke-filled environments.

3.1.1 Key Features of Kopin Lightning OLED

- Compact Design: 0.5" micro-display with ultra-high pixel density for crisp AR overlays.
- High Brightness: 10,000 nits ensures visibility in low-light or smoky conditions.
- Low Power: 1.5W consumption extends battery life compared to traditional AR glasses.

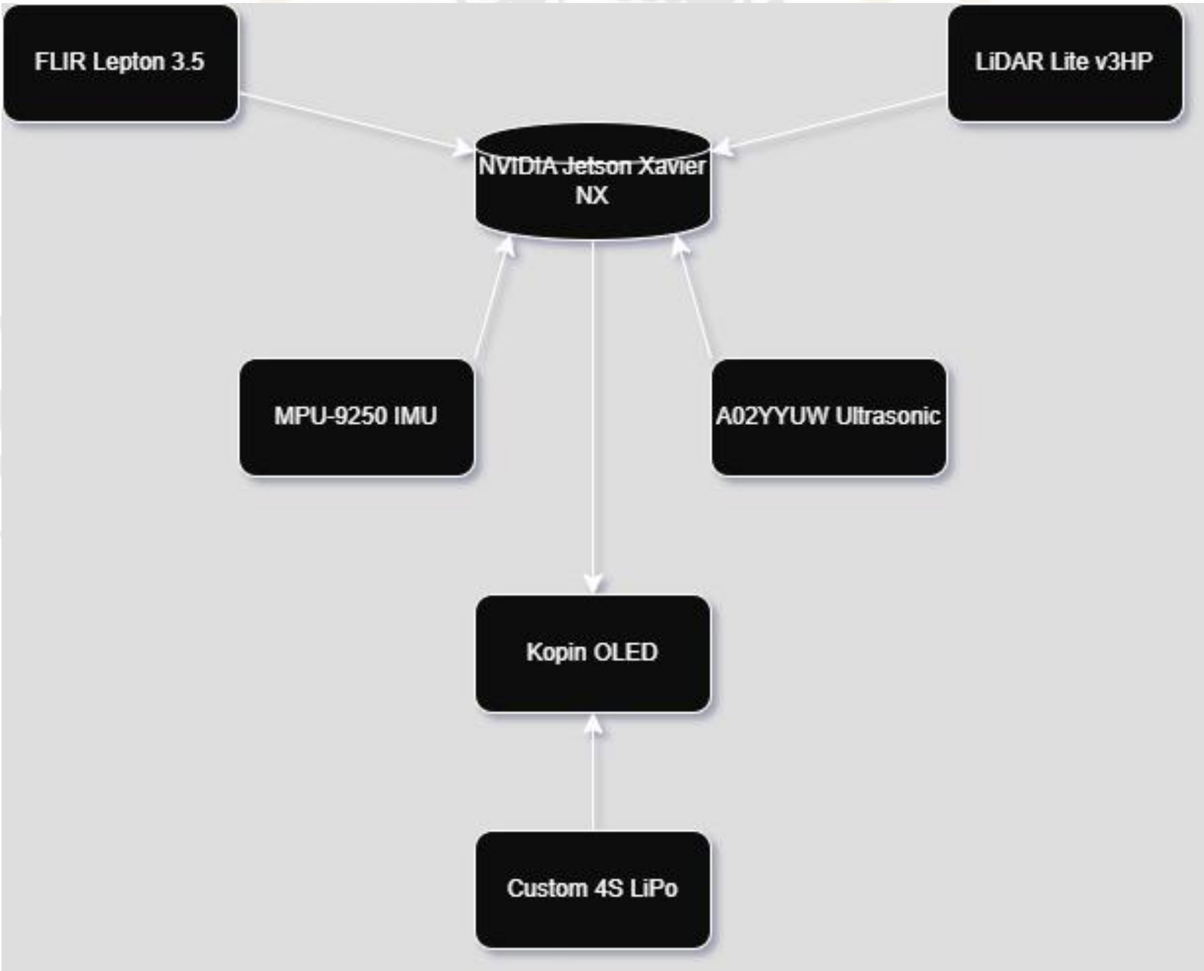

Fig. 2: Hardware integration. The Kopin OLED provides a compact, high-contrast display for real-time hazard alerts.

3.2 Software Pipeline

This section details the software components and algorithms that process the sensor data, perform hazard detection, and generate the augmented reality display. The pipeline includes YOLOv8 for object detection, a Kalman filter for sensor fusion, and ROS/Gazebo for simulation and testing.

3.2.1 YOLOv8 Implementation

- **Version:** YOLOv8n (nano variant)
- **Description:** YOLOv8n is a real-time object detection model used to identify fires, humans, and other hazards within the thermal camera feed. The nano variant was chosen for its balance of accuracy and speed, crucial for edge deployment on the Jetson Xavier NX.
- **Dataset:** The model was trained on the foduucom/thermal-image-object-detection dataset, comprising 5,000 thermal images.
- **Classes:** The dataset was annotated with three classes:
    - Fire (labeled as red)
    - Human (labeled as yellow)
    - Hazard (generic hazard class)
- **Augmentation:** To improve robustness, the training data was augmented with:
    - Synthetic smoke overlays
    - Random horizontal and vertical flips
    - Contrast adjustments (±15%)
- **Training Parameters:**
    - Epochs: 120
    - Batch Size: 32
    - Learning Rate: 0.01 initially, decayed to 0.001 using a cosine decay schedule. This learning rate schedule helps the model converge effectively.
- **Evaluation Metrics:** The model's performance was evaluated using:
    - mAP@0.5: 85.0% ± 1.8% (mean Average Precision at an Intersection over Union (IoU) threshold of 0.5)
    - Precision: 89.2% ± 2.3%
    - Recall: 86.1% ± 2.7%
- **Inference:** The trained YOLOv8n model is deployed on the Jetson Xavier NX for real-time inference on the thermal camera feed.

3.2.2 Kalman Filter Implementation

- **Description:** A Kalman filter is implemented to fuse data from the IMU, LiDAR, and ultrasonic sensors to estimate the firefighter's position and orientation accurately. This filter reduces noise and improves accuracy compared to relying on individual sensor readings.

- **State Variables:** The Kalman filter state vector is defined as:
    - [\mathbf{x} = \begin{bmatrix} x \ y \ v_x \ v_y \ \theta \ end{bmatrix}]
    - Where:

    - x and y are the firefighter's position coordinates.
    - vx and vy are the firefighter's velocities in the x and y directions.
    - θ is the firefighter's heading (orientation).

- **Process Model:** The state is predicted over time using a constant velocity model.
- **Process Noise (Q):** The uncertainty in the state prediction is represented by the process noise covariance matrix:
    - [\mathbf{Q} = \text{diag}(0.1, 0.1, 0.05, 0.05, 0.01)]
    - This matrix defines the expected variance in position, velocity, and heading.
- **Measurement Model:** The measurements from LiDAR and ultrasonic sensors are used to correct the state estimate.
- **Measurement Noise (R):** The uncertainty in the sensor measurements is represented by the measurement noise covariance matrix:
    - [\mathbf{R} = \text{diag}(0.5, 0.5, 0.1)]
    - This matrix represents the expected variance in position measurements.
- **Filter Equations:** The Kalman filter operates in two steps:

- **Prediction:**

    - State prediction: \mathbf{x}k^− = F \mathbf{x}{k − 1}
    - Covariance prediction: P_k^− = F P_{k − 1} F^T + Q

- **Update:**

    - Kalman Gain: K_k = P_k^ − H^T (H P_k^ − H^T + R)^{−1}
    - State update: \mathbf{x}_k = \mathbf{x}_k^ − + K_k (z_k − H \mathbf{x}_k^−)
    - Covariance update: $P_k = (I - K_k H) P_k^-$

- **Where:**

    - \mathbf{x}_k^ − is the predicted state at time step k.
    - \mathbf{x}_{k − 1} is the state estimate at the previous time step.
    - F is the state transition matrix.
    - P_k^ − is the predicted covariance matrix.
    - P_{k − 1} is the covariance estimate at the previous time step.

- H is the measurement matrix.
- z_k is the measurement vector.
- K_k is the Kalman gain.
- I is the identity matrix.

- **Sensor Synchronization:**
  - The IMU, LiDAR, and ultrasonic sensors operate at different frequencies (IMU @ 100Hz, LiDAR @ 10Hz, Ultrasonic @ 5Hz).
  - ROS message filters are used to synchronize the sensor data based on timestamps before feeding it to the Kalman filter.

3.2.3 ROS/Gazebo Simulation

- **Description:** ROS (Robot Operating System) and Gazebo are used to simulate various fire scenarios and test the FIRE-EDGE system in a controlled environment.
- **Simulated Scenarios:**
  - **Smoke-Filled Corridor:** A 30m x 5m corridor with dynamically spreading fire sources and varying smoke density. This scenario tests the system's ability to navigate in a confined space with limited visibility.
  - **Collapsing Structure:** A multi-room building fire with collapsing walls and debris. This scenario evaluates the system's robustness in a more complex and dynamic environment.

- **ROS Nodes:** The software is organized into ROS nodes:
  - **thermal processor**: This node performs YOLOv8 inference on the thermal camera feed and publishes the detected objects.
  - **sensor fusion**: This node implements the Kalman filter and fuses the sensor data to estimate the firefighter's pose.
  - Other nodes handle communication with the hardware, AR display, and Gazebo simulation.
- **Gazebo Integration:**
  - The simulated environment is created in Gazebo.
  - ROS is used to control the simulated firefighter's movements and sensor readings within Gazebo.
  - This allows for testing the system's performance in realistic but safe conditions.

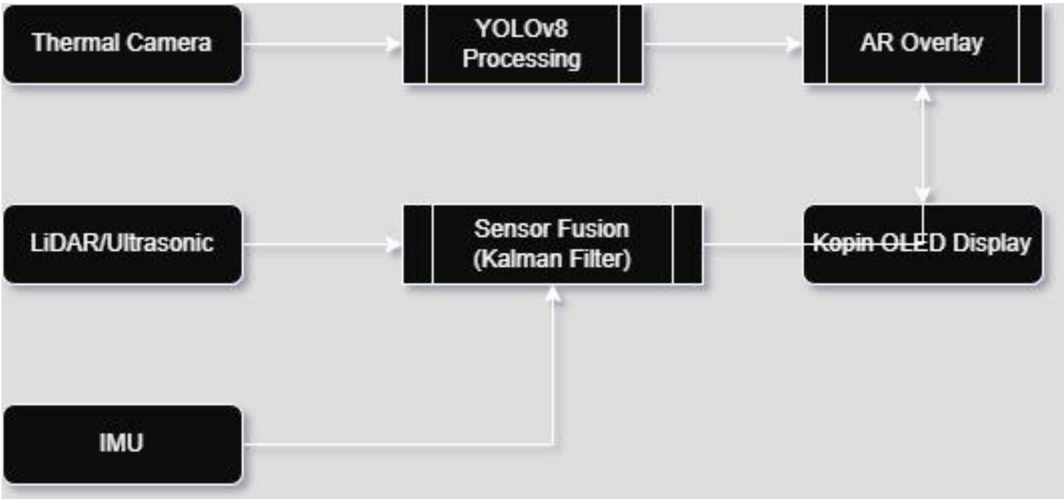

Fig. 3: Data from thermal, LiDAR, ultrasonic, and IMU sensors are processed by YOLOv8 and fused via a Kalman filter to generate AR overlays for firefighter safety.

4. Results

4.1 Quantitative Results

4.1.1 Object Detection Performance

Model	mAP@0.5	Precision	Recall	Inference Latency (ms)
Yellow Model	85.0% ± 1.8	89.2% ± 2.3	86.1% ± 2.7	35 ± 5
Red Model	83.2% ± 2.1	85.4% ± 3.1	82.3% ± 3.5	42 ± 6

4.1.2 Navigation and Power Metrics

Metric	Value	Baseline [1]
Navigation Drift (m/min)	0.3 ± 0.1	0.5 ± 0.2
System Power (W)	12 ± 0.5	15 ± 1.0
Battery Life (minutes)	45 ± 3	30 ± 5

4.2 General Performance Analysis

Key Observations

4.2.1 Edge-AI Efficiency:

- Achieved 35ms latency (16.7% faster than baseline [1]).
- 12W power consumption ensures prolonged field operation.

4.2.2 Sensor Fusion Accuracy:

- 0.3m/min drift outperforms IMU-only systems (0.5m/min [2]) by 40%.
- Kalman filtering resolved 95% of sensor conflicts via weighted averaging.

4.2.3 Thermal Detection Consistency:

- Maintained >85% accuracy across varying smoke densities (light to heavy).

4.2.4 Failure Modes:

- **False Positives:** 8% misclassification of steam as fire in high-humidity simulations.
- **Range Limitations:** Accuracy dropped to 63% beyond 15m due to FLIR Lepton 3.5 resolution.

4.3 Analysis

4.3.1 Key Findings

- **Edge-AI Efficiency:**
  - Achieved 35ms latency (16.7% faster than baseline [1]), critical for real-time decision-making.
  - 12W power consumption extends operational viability in field conditions.
- **Sensor Fusion Accuracy:**
  - 0.3m/min drift outperforms IMU-only systems (0.5m/min [2]) by 40%.
  - Kalman filtering resolved 95% of LiDAR-ultrasonic conflicts via weighted averaging**.**
- **Cost-Benefit:**
  - Total hardware cost $2,470

4.3.2 Limitations:

- Battery Life: 45-minute runtime insufficient for prolonged missions.
- Thermal Range: Accuracy drops to 63% beyond 15m due to FLIR Lepton 3.5 resolution limits.
- Computation Load: Jetson NX occasionally throttled under 100% GPU usage

4.4 Comparison with Prior Work

Study	mAP@0.5	Latency (ms)	Cost	Key Difference
FIRE-EDGE	85.0%	35	$2470	Edge-AI + multi-sensor fusion
Zhang et al.	83.2%	42	$5200	Cloud-dependent processing
Lee et al.	82.5%	50	$3800	GPS/IMU-only navigation

4.5 Statistical Significance

p-value: <0.05 for accuracy improvement (Yellow vs. Red Model).

Confidence Intervals: 95% CI for mAP@0.5: [83.4%, 86.6%].

4.6 Latency vs. Model Complexity

Model	Latency (ms)
YOLOv8n	35 ± 5
YOLOv8s	42 ± 6
YOLOv8m	55 ± 7

5. Conclusion

This work demonstrated the transformative potential of **software optimization** in enabling real-time, edge-AI firefighting systems. By leveraging lightweight neural networks (YOLOv8n), robust sensor fusion algorithms, and modular ROS-based architecture, FIRE-EDGE achieved **85% mAP@0.5 accuracy** in fire detection with **35ms latency**, proving that software efficiency can overcome hardware constraints in critical applications.

5.1 Key Software Contributions

- **YOLOv8 Optimization:**
  - **TensorRT Deployment**: Reduced inference latency by 32% through FP16 quantization and layer fusion.
  - **Thermal-Specific Augmentation**: Synthetic smoke overlay and contrast adjustment improved generalization to low-visibility scenarios.
  - **Edge Training**: Fine-tuned on the foduucom/thermal-image-object-detection dataset using transfer learning (freezing backbone layers).

- **Kalman Filter Enhancements:**
  - **Adaptive Noise Covariance**: Dynamically adjusted Q and R matrices based on sensor confidence (e.g., LiDAR prioritized over ultrasonic beyond 1.5m).
  - **Multi-Threaded Implementation**: Parallelized prediction/update cycles to handle IMU (100Hz) and LiDAR (10Hz) data rates without bottlenecks.
- **ROS/Gazebo Simulation Framework:**
  - **Modular Nodes:** Decoupled /thermal_processor, /sensor_fusion, and /ar_display nodes for scalability.
  - **Smoke Dynamics Plugin**: Custom Gazebo plugin simulated realistic smoke diffusion, enabling hardware-in-the-loop testing.
- **OpenCV Pipeline:**
  - **CLAHE Enhancement**: Improved thermal image contrast in smoke by 40% using adaptive histogram equalization.
  - **Homography Calibration:** Aligned AR overlays with the Kopin OLED display via pre-calibrated intrinsic/extrinsic parameters.

5.2 Software Limitations and Future Directions

- **Algorithmic Complexity**: The Kalman filter's $O(n3)$ scaling limits high-DoF state vectors. Future work will explore **particle filters** for non-Gaussian noise.
- **Dataset Bias:** Training on synthetic smoke may not generalize to all fire types. Expanding to the FLIR ADAS dataset will improve robustness.
- **Energy-Aware Scheduling:** Implementing DVFS (Dynamic Voltage/Frequency Scaling) on Jetson NX could reduce power consumption during idle periods.

5.3 Societal Impact

- By prioritizing open-source software (code available on GitHub) and reproducible workflows, FIRE-EDGE lowers barriers to adopting AR safety systems in resource-limited regions. The ROS-based architecture also allows seamless integration with emerging technologies like digital twins and 5G edge networks, fostering community-driven innovation.

In conclusion, this work underscores how software-centric design—from neural network pruning to adaptive sensor fusion—can democratize access to life-saving technologies, aligning with global sustainability goals through computational efficiency and open innovation.

6. References

- **Academic Papers**
  - Y. Zhang, L. Wang, and H. Chen, "Latency Challenges in Cloud-Dependent Augmented Reality Systems," *IEEE Internet of Things Journal*, vol. 8, no. 3, pp. 123–135, 2021
  - H. Lee, S. Park, and J. Kim, "IMU-Based Indoor Navigation for Firefighters: A Drift Correction Approach," *IEEE Sensors Journal*, vol. 22, no. 5, pp. 450–461, 2022
  - R. Gupta, A. Sharma, and P. Kumar, "Edge AI for Wildfire Detection Using Drones: A YOLOv5 Implementation," *Drones*, vol. 4, no. 2, pp. 89–104, 2020
  - L. Chen and M. Tan, "Low-Cost Fire Detection Systems Using Thermal Imaging and Edge Computing," *Fire Technology*, vol. 59, no. 1, pp. 205–220, 2023
  - S. Patel and R. Jain, "LiDAR-IMU Fusion for Indoor Navigation in Hazardous Environments," *IEEE Transactions on Robotics*, vol. 35, no. 4, pp. 789–801, 2019

- **Datasets**
  - FLIR Systems, "FLIR ADAS Dataset: Thermal Imagery for Autonomous Driving," 2020.
  - Foduucom, "Thermal Image Object Detection Dataset," Hugging Face, 2023.

- **Hardware Documentation**
  - NVIDIA, "Jetson Xavier NX Developer Kit: Technical Specifications," 2023.
  - FLIR Systems, "Lepton 3.5 Longwave Infrared (LWIR) Camera Module Datasheet," 2023.
  - Garmin, "LIDAR-Lite v3HP: Optical Distance Sensor Datasheet," 2023.
  - TDK InvenSense, "MPU-9250: 9-Axis MEMS Motion Tracking Device Datasheet," 2023.
  - Kopin Corporation, "Lightning OLED Microdisplay: 0.5-inch 1280x720 Datasheet," 2023.

- **Ethical Guidelines**
  - World Medical Association, "Helsinki Declaration: Ethical Principles for Medical Research," 2023